


File System Management

Practice Exercises – Solutions

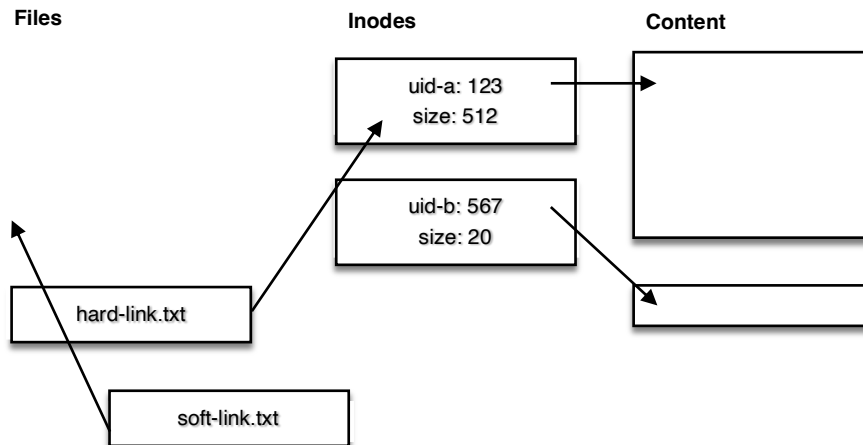
COMP-310/ECSE-427 Winter 23

Question 1

A.



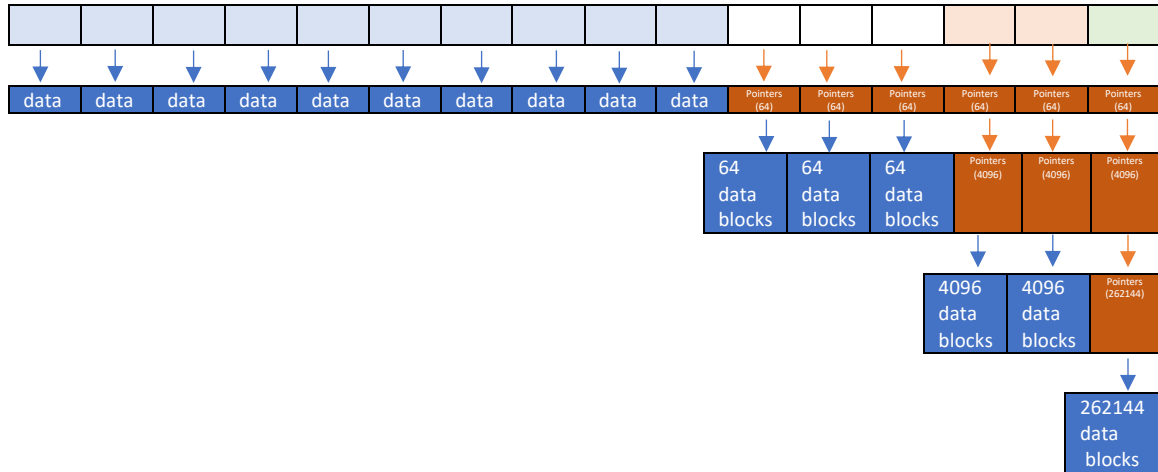
B. `soft-link.txt` will be a **dangling reference**, while `hard-link.txt` will **still point to the inode** of `a.txt`.



Question 2

A. $10 * 512B + 3 * 64 * 512B + 2 * 64^2 * 512B + 64^3 * 512B = 138,515,456 B \approx 13 MB$

B.



In this case, the user cannot open a 1GB file, because the maximum file size is 13MB. No disk accesses occur because the requested file is too large.

- C. Since we have 10 direct pointers, it is likely that the typical file size is < 5KB because files < 5KB are the fastest to access. Good file systems design needs to optimize for the common scenarios, while supporting the uncommon cases (e.g., files larger than 5KB).

Question 3

Max file size = $PD * S + PI * P * S + PDI * P^2 * S = S(PD + PI * P + PDI * P^2)$ Bytes

s

Question 4

Out of SCAN (moving up), C-LOOK (moving up, serving requests on the way up), and SSTF, which disk scheduling policy minimizes the total hard disk arm motion?

→ SSTF

Question 5**Linked Allocation**

- Add block in the beginning: 1 I/O
- Add block in the middle: 102 I/Os
- Add block in the end: 202 I/Os

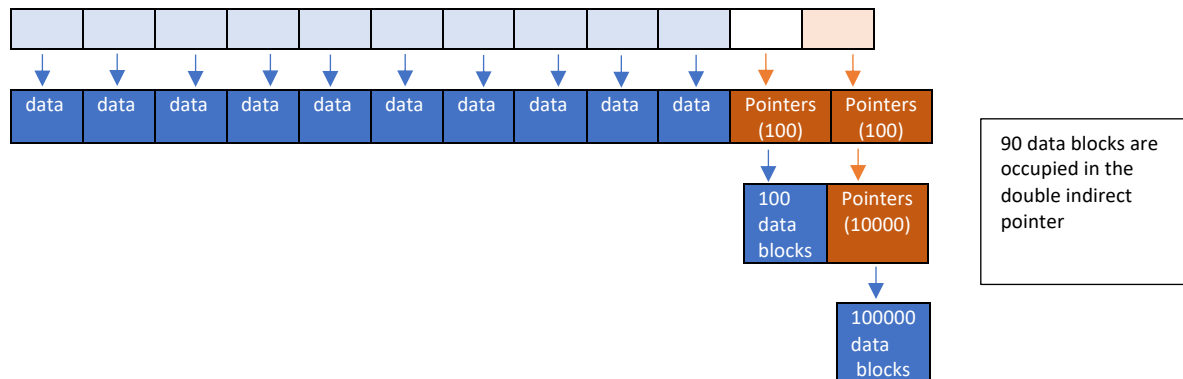
FAT

- Add block in the beginning: 1 I/O
- Add block in the middle: 1 I/O

c) Add block in the end: 1 I/O

Indexed allocation with 10 direct pointers, 1 indirect pointer, and 1 double-indirect pointer, where each pointer block can hold 100 pointers

The Inode is in memory, so the first layer of pointers is cached.



- Add block in the beginning: 1 I/O to write block to disk + 1 + 2*2 I/Os to modify the indirect pointer and the second level of the double-indirect pointer = 6 I/Os
- Add block in the middle: 1 I/O to write block to disk + 1 + 2*2 I/Os to modify the indirect pointer and the second level of the double-indirect pointer = 6 I/Os
- Add block in the end: : 1 I/O to write block to disk + 1 + 2 I/O to modify the second level of the double-indirect pointer = 4 I/Os

Question 6

```
fd = Open(bigfile)
for (i = 0; i < 100; i++) {
    R = randomfilelocation()
    Seek( fd, R)
    if (R<2^30-10)
        Read( fd, buffer, 10)
    else
        Read(fd, buffer, 2^30-R) //if we're close to eof we cannot read 10 Bytes
    printf(buffer)
}
Close(fd)
```

Question 7

- The main advantage of RAID-5 over RAID-4 is more efficient handling of random writes, as RAID-5 does not suffer from the small writes problem. The reason is that the parity blocks are distributed across the entire disk array in RAID-5, as opposed to RAID-4

where the parity blocks reside in a single disk. The parity disk becomes the bottleneck in RAID-4.

- B. The minimum number of disks to configure a RAID-5 array is 3. RAID 5 can sustain the loss of a single drive. In the event of a drive failure, data from the failed drive is reconstructed from parity striped across the remaining drives. Therefore, RAID-5 cannot be configured with just 2 drives, as the equivalent capacity of one of the drives needs to be used for the parity blocks.

Question 8

- A. In memory: LFS maintains pointers to the Imap pieces, which, when accessed can retrieve the full Imap from disk.
- B. On disk: LFS maintains the actual Imap, spread across different segments, as well as the data that is referenced by the Imap. Optional: LFS can maintain a checkpointing region for the Imap on disk, but this is an optimization. It is possible to decide what blocks are old without using the checkpointing region.
- C. To figure out whether a block B is logically overwritten by a later write to the same block, LFS does the following:
 - Scan the entire Imap and check whether B is referenced by any of the Inodes.
 - If B is referenced by an Inode, it means that B has not been logically replaced by a newer version.
 - If B is not referenced, it means B has been logically replaced and should be garbage collected.